

Tutorial on Bayesian Optimization

Johannes Kirschner and Mojmír Mutný

February 26th, 2019

ICFA ML Workshop, PSI



Motivating Application: Parameter Tuning of Accelerator



Maximize (photon) signal, minimize losses, ...

[McIntire et al., 2016, Kirschner et al., 2019]


Motivating Application: Experimental Design



Optimize design parameters, e.g. nano materials, molecules,...

[Schneider et al., 2018, Romero et al., 2013]

Motivating Application: Fitting Physical Models



Optimize model parameters to fit observational data
[Ekström et al., 2019]

Optimizing Black-Box Functions

Maximize a Black box function:

$$\mathcal{X} \longrightarrow f \rightsquigarrow f(x) + \epsilon$$

Optimizing Black-Box Functions

Maximize a Black box function:

$$\mathcal{X} \longrightarrow f \rightsquigarrow f(x) + \epsilon$$

- ▷ Parameter space $\mathcal{X} \subset \mathbb{R}^d$, can also be combinatorial
- ▷ Little assumptions on f : non-convex, multiple local optima, ...
- ▷ No analytical formula for f
- ▷ No access to gradients

Optimizing Black-Box Functions

Maximize a Black box function:

$$\mathcal{X} \longrightarrow f \rightsquigarrow f(x) + \epsilon$$

- ▷ Parameter space $\mathcal{X} \subset \mathbb{R}^d$, can also be combinatorial
- ▷ Little assumptions on f : non-convex, multiple local optima, ...
- ▷ No analytical formula for f
- ▷ No access to gradients

Only get (noisy) evaluations $y = f(x) + \epsilon$

- ▷ Evaluations of f are ‘expensive’

Bayesian Optimization: Overview

Prior data set: \mathcal{D}_0

Bayesian Optimization: Overview

Prior data set: \mathcal{D}_0

For each step $t = 1, 2, 3, \dots, T$,

Step 1: Build *probabilistic model* \hat{f}_t of the objective using data \mathcal{D}_{t-1}

▷ Gaussian process regression (**Part I**)

Bayesian Optimization: Overview

Prior data set: \mathcal{D}_0

For each step $t = 1, 2, 3, \dots, T$,

Step 1: Build *probabilistic model* \hat{f}_t of the objective using data \mathcal{D}_{t-1}

▷ Gaussian process regression (**Part I**)

Step 2: Reduce model *uncertainty about maximizers*

▷ Search guided by acquisition function $x_t = \arg \max_{x \in \mathcal{X}} \alpha(x|\hat{f}_t)$ (**Part II**)

Bayesian Optimization: Overview

Prior data set: \mathcal{D}_0

For each step $t = 1, 2, 3, \dots, T$,

Step 1: Build *probabilistic model* \hat{f}_t of the objective using data \mathcal{D}_{t-1}

▷ Gaussian process regression (**Part I**)

Step 2: Reduce model *uncertainty about maximizers*

▷ Search guided by acquisition function $x_t = \arg \max_{x \in \mathcal{X}} \alpha(x|\hat{f}_t)$ (**Part II**)

Step 3: Observe (noisy) measurement $y_t = f(x_t) + \epsilon$

▷ Augment data $\mathcal{D}_t = \mathcal{D}_{t-1} \cup \{(x_t, y_t)\}$

Bayesian Optimization: Overview

Prior data set: \mathcal{D}_0

For each step $t = 1, 2, 3, \dots, T$,

Step 1: Build *probabilistic model* \hat{f}_t of the objective using data \mathcal{D}_{t-1}

▷ Gaussian process regression (**Part I**)

Step 2: Reduce model *uncertainty about maximizers*

▷ Search guided by acquisition function $x_t = \arg \max_{x \in \mathcal{X}} \alpha(x|\hat{f}_t)$ (**Part II**)

Step 3: Observe (noisy) measurement $y_t = f(x_t) + \epsilon$

▷ Augment data $\mathcal{D}_t = \mathcal{D}_{t-1} \cup \{(x_t, y_t)\}$

At final time T : *Use model to find best predicted setting.*

Part I: Gaussian Process Regression

Bayesian Statistics

Prior: Distribution $\mathcal{P}(f)$ over f

- ▷ “prior belief”

Bayesian Statistics

Prior: Distribution $\mathcal{P}(f)$ over f

- ▷ “prior belief”

Data likelihood: $\mathcal{P}(D_t|f)$

- ▷ e.g. $y \sim f(x) + \mathcal{N}(0, 1)$

Bayesian Statistics

Prior: Distribution $\mathcal{P}(f)$ over f

- ▷ “prior belief”


Data likelihood: $\mathcal{P}(D_t|f)$

- ▷ e.g. $y \sim f(x) + \mathcal{N}(0, 1)$


Posterior distribution: $\mathcal{P}(f|D_t) = \frac{\mathcal{P}(D_t|f)\mathcal{P}(f)}{\mathcal{P}(D_t)}$

- ▷ Bayes’ theorem
- ▷ The posterior distribution captures our belief in f after seeing the data.


Gaussian Processes



Normal dist.
(1-D Gaussian)




Multivariate normal
(n-D Gaussian)



Gaussian process
(∞ -D Gaussian)


- ▷ **Gaussian process (GP)** = normal distribution over *functions*

Gaussian Processes



- ▷ **Gaussian process (GP)** = normal distribution over *functions*
- ▷ Finite marginals $f(x_1), \dots, f(x_n)$ are multivariate Gaussians

Gaussian Processes



- ▷ **Gaussian process (GP)** = normal distribution over *functions*
- ▷ Finite marginals $f(x_1), \dots, f(x_n)$ are multivariate Gaussians
- ▷ Parameterized by covariance function (kernel) $k(x, x') = \text{Cov}(f(x), f(x'))$

Gaussian Process on Finite Domain

Finite domain: $\mathcal{X} = \{x_1, \dots, x_n\}$

Gaussian Process on Finite Domain

Finite domain: $\mathcal{X} = \{x_1, \dots, x_n\}$

Definition: f is a Gaussian process with mean $\mu(x)$ and kernel $k(x, x')$

▷ if $f(x_1), \dots, f(x_n)$ is multivariate normal $\mathcal{N}(m, K)$ with

Gaussian Process on Finite Domain

Finite domain: $\mathcal{X} = \{x_1, \dots, x_n\}$

Definition: f is a Gaussian process with mean $\mu(x)$ and kernel $k(x, x')$

- ▷ if $f(x_1), \dots, f(x_n)$ is multivariate normal $\mathcal{N}(m, K)$ with
- ▷ mean $m = [\mu(x_1), \dots, \mu(x_n)]$,

Gaussian Process on Finite Domain

Finite domain: $\mathcal{X} = \{x_1, \dots, x_n\}$

Definition: f is a Gaussian process with mean $\mu(x)$ and kernel $k(x, x')$

- ▷ if $f(x_1), \dots, f(x_n)$ is multivariate normal $\mathcal{N}(m, K)$ with
 - ▷ mean $m = [\mu(x_1), \dots, \mu(x_n)]$,
 - ▷ covariance $K = [k(x_i, x_j)]_{i,j=1,\dots,n}$.

Gaussian Process on Finite Domain

Finite domain: $\mathcal{X} = \{x_1, \dots, x_n\}$

Definition: f is a Gaussian process with mean $\mu(x)$ and kernel $k(x, x')$

- ▷ if $f(x_1), \dots, f(x_n)$ is multivariate normal $\mathcal{N}(m, K)$ with
 - ▷ mean $m = [\mu(x_1), \dots, \mu(x_n)]$,
 - ▷ covariance $K = [k(x_i, x_j)]_{i,j=1,\dots,n}$.

Denote $f \sim GP(m, k)$.

Gaussian Process on Continuous Domain

Continuous domain: $\mathcal{X} \subset \mathbb{R}^d$

Gaussian Process on Continuous Domain

Continuous domain: $\mathcal{X} \subset \mathbb{R}^d$

Definition: f is a Gaussian process with mean $\mu(x)$ and kernel $k(x, x')$ if

- ▷ for **any finite subset** $\{x_1, \dots, x_n\} \subset \mathcal{X}$,

Gaussian Process on Continuous Domain

Continuous domain: $\mathcal{X} \subset \mathbb{R}^d$

Definition: f is a Gaussian process with mean $\mu(x)$ and kernel $k(x, x')$ if

- ▷ for **any finite subset** $\{x_1, \dots, x_n\} \subset \mathcal{X}$,
- ▷ $f(x_1), \dots, f(x_n)$ is multivariate normal $\mathcal{N}(m, K)$ with
- ▷ mean $m = [\mu(x_1), \dots, \mu(x_n)]$,
- ▷ covariance $K = [k(x_i, x_j)]_{i,j=1,\dots,n}$.

Gaussian Process on Continuous Domain


Continuous domain: $\mathcal{X} \subset \mathbb{R}^d$

Definition: f is a Gaussian process with mean $\mu(x)$ and kernel $k(x, x')$ if


- ▷ for **any finite subset** $\{x_1, \dots, x_n\} \subset \mathcal{X}$,
- ▷ $f(x_1), \dots, f(x_n)$ is multivariate normal $\mathcal{N}(m, K)$ with
- ▷ mean $m = [\mu(x_1), \dots, \mu(x_n)]$,
- ▷ covariance $K = [k(x_i, x_j)]_{i,j=1,\dots,n}$.

In practice we always evaluate/sample the GP on finite (grid) domains.


Samples from a Gaussian Process



Samples from a Gaussian Process



Samples from a Gaussian Process



Gaussian Process Regression

Prior: GP prior $\mathcal{P}(f) = GP(\mu, k)$ over f

- ▷ “prior belief” with prior mean μ and kernel k

Gaussian Process Regression

Prior: GP prior $\mathcal{P}(f) = GP(\mu, k)$ over f

- ▷ “prior belief” with prior mean μ and kernel k

Gaussian likelihood: iid Gaussian noise:

- ▷ $\mathcal{P}(\{y_1, \dots, y_m\} | f(x_1), \dots, f(x_n)) = \prod_i \mathcal{N}(f(x_i), \rho^2)$
- ▷ e.g. $y \sim f(x) + \mathcal{N}(0, \rho^2)$

Gaussian Process Regression

Prior: GP prior $\mathcal{P}(f) = GP(\mu, k)$ over f

- ▷ “prior belief” with prior mean μ and kernel k

Gaussian likelihood: iid Gaussian noise:

- ▷ $\mathcal{P}(\{y_1, \dots, y_m\} | f(x_1), \dots, f(x_n)) = \prod_i \mathcal{N}(f(x_i), \rho^2)$
- ▷ e.g. $y \sim f(x) + \mathcal{N}(0, \rho^2)$


Posterior distribution: $\mathcal{P}(f | D_t) = GP(\mu_n, k_n)$

- ▷ Posterior distributions is again a GP!
- ▷ Closed form updates exist.
- ▷ Excellent book (free pdf): [Rasmussen, 2004, Chapter 2]

Marginals

Posterior distribution: $\mathcal{P}(f|D_t) = GP(\mu_n, k_n)$


- ▷ *Remember:* Finite marginals are Gaussians!
- ▷ Marginal posterior distribution at any point x is $\mathcal{N}(\mu_n(x), k_n(x, x))$



Marginals

Posterior distribution: $\mathcal{P}(f|D_t) = GP(\mu_n, k_n)$

- ▷ Remember: Finite marginals are Gaussians!
- ▷ Marginal posterior distribution at any point x is $\mathcal{N}(\mu_n(x), k_n(x, x))$



Posterior variance $\sigma_n(x)^2 = k_n(x, x)$ quantifies uncertainty

Kernel Functions

Kernel k needs to satisfy some technical assumptions:

- ▷ symmetric
- ▷ positive semidefinite.


Kernel Functions

Kernel k needs to satisfy some technical assumptions:

- ▷ symmetric
- ▷ positive semidefinite.

Kernels are similarity measures between points and encodes smoothness.


Kernel Functions: Squared Exponential (RBF)



Squared exponential kernel: $k(x, x') = \exp(-\|x - x'\|^2 / l^2)$

- ▷ l is called lengthscale (or bandwidth)


Kernel Functions: Exponential



Exponential kernel: $k(x, x') = \exp(-\|x - x'\|^2/l^2)$

- ▷ l is called lengthscale (or bandwidth)


Kernel Functions: Matern



$$\text{Matern32 kernel: } k(x, x') = \left(1 + \frac{\sqrt{3}\|x - x'\|}{l}\right) \exp\left(-\frac{\sqrt{3}\|x - x'\|}{l}\right)$$

- ▷ l is called lengthscale (or bandwidth)
- ▷ Matern52, etc: Family of kernels with increasing smoothness

Kernel Functions: Linear



Linear kernel: $k(x, x') = x^\top x'$

- ▷ Recovers (Bayesian) linear regression!

Feature kernel: $k(x, x') = \Phi(x)^\top \Phi(x')$

- ▷ E.g. polynomials $\Phi(x) = [1, x, x^2]$

Kernel Parameters I

Noise variance

- ▷ Easy to measure
- ▷ Slightly larger value increases robustness

Kernel Parameters I

Noise variance

- ▷ Easy to measure
- ▷ Slightly larger value increases robustness

Kernel

- ▷ Smoothness of function
- ▷ RBF smooth functions
- ▷ Matern32, Matern52, less smooth, often work well in practice
- ▷ Can also combine kernels, e.g. RBF + 5·Matern32
- ▷ Each kernel has its own hyper-parameters

Kernel Parameters II

Normalizes objective (y-values)

Prior variance

- ▷ Expected range of objective values
- ▷ Keep fixed (to 1) and normalize data

Kernel Parameters II

Normalizes objective (y-values)

Prior variance

- ▷ Expected range of objective values
- ▷ Keep fixed (to 1) and normalize data

Normalizes domain (x-values)

Lengthscale

- ▷ Smoothness of function
- ▷ If too large, might not model the objective well
- ▷ Can pick different lengthscales for different dimensions (ARD)
- ▷ Normalizes the domain

How to choose parameters?

Try and error

- ▷ Parameters usually more intuitive to tune

How to choose parameters?

Try and error

- ▷ Parameters usually more intuitive to tune

Point estimates

- ▷ Maximum a posteriori estimation: $\theta^* = \arg \max_{\theta} \mathcal{P}(D_t | \theta) \mathcal{P}(\theta)$
- ▷ Requires ‘representative’ initial data
- ▷ Might not work well with data collected while optimizing

How to choose parameters?

Try and error

- ▷ Parameters usually more intuitive to tune

Point estimates

- ▷ Maximum a posteriori estimation: $\theta^* = \arg \max_{\theta} \mathcal{P}(D_t | \theta) \mathcal{P}(\theta)$
- ▷ Requires ‘representative’ initial data
- ▷ Might not work well with data collected while optimizing

Bayesian approach

- ▷ Define ‘reasonable’ prior distribution $\mathcal{P}(\theta)$ over θ
- ▷ Marginalize predictions over posterior $\mathcal{P}(\theta | D_t)$
- ▷ More expensive to compute, no closed form
- ▷ Eliminates hyperparameters

Notebook Session: GP Regression using GPy

Part II: Bayesian Optimization

Optimization - recap

- ▷ Assume function $f(x)$ where $x \in \mathcal{X}$.

Optimization - recap


- ▷ Assume function $f(x)$ where $x \in \mathcal{X}$.
- ▷ Noisy zero-order oracle $\iff y = f(x) + \epsilon$

Optimization - recap

- ▷ Assume function $f(x)$ where $x \in \mathcal{X}$.
- ▷ Noisy zero-order oracle $\iff y = f(x) + \epsilon$
- ▷ Grid approach fails:


Optimization - recap

- ▷ Assume function $f(x)$ where $x \in \mathcal{X}$.
- ▷ Noisy zero-order oracle $\iff y = f(x) + \epsilon$
- ▷ Grid approach fails:
 - ▷ due to noise



Optimization - recap


- ▷ Assume function $f(x)$ where $x \in \mathcal{X}$.
- ▷ Noisy zero-order oracle $\iff y = f(x) + \epsilon$
- ▷ Grid approach fails:
 - ▷ due to noise



Optimization - recap

- ▷ Assume function $f(x)$ where $x \in \mathcal{X}$.
- ▷ Noisy zero-order oracle $\iff y = f(x) + \epsilon$
- ▷ Grid approach fails:

- ▷ due to noise




- ▷ due to efficiency [to come]

Optimization - recap


- ▷ Assume function $f(x)$ where $x \in \mathcal{X}$.
- ▷ Noisy zero-order oracle $\iff y = f(x) + \epsilon$
- ▷ Grid approach fails:

- ▷ due to noise




- ▷ due to efficiency [to come]


GPs: Avoid not necessary evaluations




GPs: Avoid not necessary evaluations




GPs: Avoid not necessary evaluations



GPs: Avoid not necessary evaluations



GPs: Avoid not necessary evaluations



Bayesian Optimization: Overview

Prior data set: \mathcal{D}_0

Bayesian Optimization: Overview

Prior data set: \mathcal{D}_0

For each step $t = 1, 2, 3, \dots, T$,

Step 1: Build *probabilistic model* \hat{f}_t of the objective using data \mathcal{D}_{t-1}

▷ Gaussian process regression (**Part I**)

Bayesian Optimization: Overview

Prior data set: \mathcal{D}_0

For each step $t = 1, 2, 3, \dots, T$,

Step 1: Build *probabilistic model* \hat{f}_t of the objective using data \mathcal{D}_{t-1}

▷ Gaussian process regression (**Part I**)

Step 2: Reduce model *uncertainty about maximizers*

▷ Search guided by acquisition function $x_t = \arg \max_{x \in \mathcal{X}} \alpha(x|\hat{f}_t)$ (**Part II**)

Bayesian Optimization: Overview

Prior data set: \mathcal{D}_0

For each step $t = 1, 2, 3, \dots, T$,

Step 1: Build *probabilistic model* \hat{f}_t of the objective using data \mathcal{D}_{t-1}

▷ Gaussian process regression (**Part I**)

Step 2: Reduce model *uncertainty about maximizers*

▷ Search guided by acquisition function $x_t = \arg \max_{x \in \mathcal{X}} \alpha(x|\hat{f}_t)$ (**Part II**)

Step 3: Observe (noisy) measurement $y_t = f(x_t) + \epsilon$

▷ Augment data $\mathcal{D}_t = \mathcal{D}_{t-1} \cup \{(x_t, y_t)\}$

Bayesian Optimization: Overview

Prior data set: \mathcal{D}_0

For each step $t = 1, 2, 3, \dots, T$,

Step 1: Build *probabilistic model* \hat{f}_t of the objective using data \mathcal{D}_{t-1}

▷ Gaussian process regression (**Part I**)

Step 2: Reduce model *uncertainty about maximizers*

▷ Search guided by acquisition function $x_t = \arg \max_{x \in \mathcal{X}} \alpha(x|\hat{f}_t)$ (**Part II**)

Step 3: Observe (noisy) measurement $y_t = f(x_t) + \epsilon$

▷ Augment data $\mathcal{D}_t = \mathcal{D}_{t-1} \cup \{(x_t, y_t)\}$

At final time T : *Use model to find best predicted setting.*

Upper Confidence Bound (UCB)

- ▷ μ_t ... *posterior mean* after seeing t points

Upper Confidence Bound (UCB)

- ▷ μ_t ... *posterior mean* after seeing t points
- ▷ σ_t ... *posterior standard deviation* after seeing t points

Upper Confidence Bound (UCB)

- ▷ μ_t ... *posterior mean* after seeing t points
- ▷ σ_t ... *posterior standard deviation* after seeing t points
- ▷ $\beta \in \mathbb{R}$ real parameter trading *exploration and exploitation* [see later]

Upper Confidence Bound (UCB)

- ▷ μ_t ... *posterior mean* after seeing t points
- ▷ σ_t ... *posterior standard deviation* after seeing t points
- ▷ $\beta \in \mathbb{R}$ real parameter trading *exploration and exploitation* [see later]

$$\alpha_t(x) = \mu_t(x) + \beta\sigma_t(x)$$

- ▷ How to optimize $\alpha_t(x)$?
 - ▷ discretize search space \mathcal{X}

Upper Confidence Bound (UCB)

- ▷ μ_t ... *posterior mean* after seeing t points
- ▷ σ_t ... *posterior standard deviation* after seeing t points
- ▷ $\beta \in \mathbb{R}$ real parameter trading *exploration and exploitation* [see later]

$$\alpha_t(x) = \mu_t(x) + \beta\sigma_t(x)$$


- ▷ How to optimize $\alpha_t(x)$?
 - ▷ discretize search space \mathcal{X}
 - ▷ first-order heuristics

Upper Confidence Bound (UCB)


- ▷ μ_t ... *posterior mean* after seeing t points
- ▷ σ_t ... *posterior standard deviation* after seeing t points
- ▷ $\beta \in \mathbb{R}$ real parameter trading *exploration and exploitation* [see later]

$$\alpha_t(x) = \mu_t(x) + \beta\sigma_t(x)$$


- ▷ How to optimize $\alpha_t(x)$?
 - ▷ discretize search space \mathcal{X}
 - ▷ first-order heuristics




UCB: Example




UCB: Example




UCB: Example




UCB: Example




UCB: Example




UCB: Example



UCB: Example



UCB: Example



Consequence of β

- ▷ What is β ?

Consequence of β

- ▷ What is β ?
- ▷ β trades exploration and exploitation

Consequence of β

- ▷ What is β ?
- ▷ β trades exploration and exploitation
- ▷ Theoretical value that ensure global convergence (right model assumption):

$$\beta_t = 2 \log(\gamma_t + 1)$$

Consequence of β

- ▷ What is β ?
- ▷ β trades exploration and exploitation
- ▷ Theoretical value that ensure global convergence (right model assumption):

$$\beta_t = 2 \log(\gamma_t + 1)$$

where γ_t is maximum information gain, for RBF kernel $\gamma_t = C \log(T)^{d+1}$

Consequence of β

- ▷ What is β ?
- ▷ β trades exploration and exploitation
- ▷ Theoretical value that ensure global convergence (right model assumption):

$$\beta_t = 2 \log(\gamma_t + 1)$$

where γ_t is maximum information gain, for RBF kernel $\gamma_t = C \log(T)^{d+1}$

- ▷ (Very common) *heuristic* approach: $\beta \approx 2$.

Consequence of β

- ▷ What is β ?
- ▷ β trades exploration and exploitation
- ▷ Theoretical value that ensure global convergence (right model assumption):

$$\beta_t = 2 \log(\gamma_t + 1)$$

where γ_t is maximum information gain, for RBF kernel $\gamma_t = C \log(T)^{d+1}$

- ▷ (Very common) *heuristic* approach: $\beta \approx 2$.
- ▷ β too small \implies gets stuck/hill climbing

Consequence of β

- ▷ What is β ?
- ▷ β trades exploration and exploitation
- ▷ Theoretical value that ensure global convergence (right model assumption):


$$\beta_t = 2 \log(\gamma_t + 1)$$

where γ_t is maximum information gain, for RBF kernel $\gamma_t = C \log(T)^{d+1}$

- ▷ (Very common) *heuristic* approach: $\beta \approx 2$.
- ▷ β too small \implies gets stuck/hill climbing
- ▷ β too high \implies incremental grid search


Consequence of β II

▷ Hill Climbing - β small




Consequence of β II

▷ Hill Climbing - β small




Consequence of β II

▷ Hill Climbing - β small




Consequence of β II

▷ Hill Climbing - β small




Consequence of β II

▷ Hill Climbing - β small




Consequence of β II

▷ Hill Climbing - β small




Consequence of β II

▷ Hill Climbing - β small




Consequence of β II

▷ Hill Climbing - β small




Consequence of β II

▷ Hill Climbing - β small




Consequence of β II

▷ Hill Climbing - β small




Consequence of β II

- ▷ Hill Climbing - β small




Consequence of β II

▷ Hill Climbing - β small




▷ Sequential Grid - β large




Consequence of β II

▷ Hill Climbing - β small




▷ Sequential Grid - β large




Consequence of β II

▷ Hill Climbing - β small




▷ Sequential Grid - β large




Consequence of β II

▷ Hill Climbing - β small




▷ Sequential Grid - β large




Consequence of β II

▷ Hill Climbing - β small




▷ Sequential Grid - β large




Consequence of β II

▷ Hill Climbing - β small




▷ Sequential Grid - β large




Consequence of β II

▷ Hill Climbing - β small



▷ Sequential Grid - β large



Other acquisition function

- ▷ *Thompson sampling*

Other acquisition function

- ▷ *Thompson sampling*
 - ▷ Sample a path $s \sim \text{GP}(\mu_t, \sigma_t)$

Other acquisition function

- ▷ *Thompson sampling*
 - ▷ Sample a path $s \sim \text{GP}(\mu_t, \sigma_t)$
 - ▷ Acquisition $\alpha_t(x) = s(x)$

Other acquisition function

- ▷ *Thompson sampling*
 - ▷ Sample a path $s \sim \text{GP}(\mu_t, \sigma_t)$
 - ▷ Acquisition $\alpha_t(x) = s(x)$
 - ▷ Empirically works often better

Other acquisition function

- ▷ *Thompson sampling*
 - ▷ Sample a path $s \sim \text{GP}(\mu_t, \sigma_t)$
 - ▷ Acquisition $\alpha_t(x) = s(x)$
 - ▷ Empirically works often better
- ▷ *Expected Improvement* [Mockus, 1982]

Other acquisition function

- ▷ *Thompson sampling*
 - ▷ Sample a path $s \sim \text{GP}(\mu_t, \sigma_t)$
 - ▷ Acquisition $\alpha_t(x) = s(x)$
 - ▷ Empirically works often better
- ▷ *Expected Improvement* [Mockus, 1982]
 - ▷ $\mu_t(x^+)$ is the best mean estimate

Other acquisition function

- ▷ *Thompson sampling*
 - ▷ Sample a path $s \sim \text{GP}(\mu_t, \sigma_t)$
 - ▷ Acquisition $\alpha_t(x) = s(x)$
 - ▷ Empirically works often better
- ▷ *Expected Improvement* [Mockus, 1982]
 - ▷ $\mu_t(x^+)$ is the best mean estimate
 - ▷ $\alpha_t(x) = \mathbb{E}[\max(0, f(x) - f(x^+)) | \mathcal{D}_t]$

Other acquisition function

- ▷ *Thompson sampling*
 - ▷ Sample a path $s \sim \text{GP}(\mu_t, \sigma_t)$
 - ▷ Acquisition $\alpha_t(x) = s(x)$
 - ▷ Empirically works often better
- ▷ *Expected Improvement* [Mockus, 1982]
 - ▷ $\mu_t(x^+)$ is the best mean estimate
 - ▷ $\alpha_t(x) = \mathbb{E}[\max(0, f(x) - f(x^+)) | \mathcal{D}_t]$
 - ▷ Analytical solution: $\alpha_t(x) = (\mu_t(x) - \mu(x^+))\Phi(Z) + \sigma(x)\phi(Z)$ where,
 $Z = \frac{\mu_t - \mu(x^+)}{\sigma_t(x)}$ and Φ, ϕ are cdf and pdf of standard normal.

Other acquisition function

- ▷ *Thompson sampling*
 - ▷ Sample a path $s \sim \text{GP}(\mu_t, \sigma_t)$
 - ▷ Acquisition $\alpha_t(x) = s(x)$
 - ▷ Empirically works often better
- ▷ *Expected Improvement* [Mockus, 1982]
 - ▷ $\mu_t(x^+)$ is the best mean estimate
 - ▷ $\alpha_t(x) = \mathbb{E}[\max(0, f(x) - f(x^+)) | \mathcal{D}_t]$
 - ▷ Analytical solution: $\alpha_t(x) = (\mu_t(x) - \mu(x^+))\Phi(Z) + \sigma(x)\phi(Z)$ where,
 $Z = \frac{\mu_t - \mu(x^+)}{\sigma_t(x)}$ and Φ, ϕ are cdf and pdf of standard normal.

Curse of dimensionality

- ▷ How do we apply this to multiple dimensions?

Curse of dimensionality

- ▷ How do we apply this to multiple dimensions?
- ▷ Naturally, $\alpha_t(x)$ can be defined in any $\mathcal{X} \subset \mathbb{R}^d$

Curse of dimensionality

- ▷ How do we apply this to multiple dimensions?
- ▷ Naturally, $\alpha_t(x)$ can be defined in any $\mathcal{X} \subset \mathbb{R}^d$
- ▷ Practically, $\alpha_t(x)$ cannot be optimized using a grid optimizer.

Curse of dimensionality

- ▷ How do we apply this to multiple dimensions?
- ▷ Naturally, $\alpha_t(x)$ can be defined in any $\mathcal{X} \subset \mathbb{R}^d$
- ▷ Practically, $\alpha_t(x)$ cannot be optimized using a grid optimizer.
- ▷ The size of the grid grows n^d and computational needs grow as $(n^d)^3$, where n number of grid points in 1D.

Curse of dimensionality

- ▷ How do we apply this to multiple dimensions?
- ▷ Naturally, $\alpha_t(x)$ can be defined in any $\mathcal{X} \subset \mathbb{R}^d$
- ▷ Practically, $\alpha_t(x)$ cannot be optimized using a grid optimizer.
- ▷ The size of the grid grows n^d and computational needs grow as $(n^d)^3$, where n number of grid points in 1D.
- ▷ One can use a first-order heuristic to optimize the acquisition function locally.

Curse of dimensionality

- ▷ How do we apply this to multiple dimensions?
- ▷ Naturally, $\alpha_t(x)$ can be defined in any $\mathcal{X} \subset \mathbb{R}^d$
- ▷ Practically, $\alpha_t(x)$ cannot be optimized using a grid optimizer.
- ▷ The size of the grid grows n^d and computational needs grow as $(n^d)^3$, where n number of grid points in 1D.
- ▷ One can use a first-order heuristic to optimize the acquisition function locally.
- ▷ More advanced methods: Look [Mutný and Krause, 2018] or visit:

Talk of Johannes tomorrow on BO for SwissFEL.

Curse of dimensionality

- ▷ How do we apply this to multiple dimensions?
- ▷ Naturally, $\alpha_t(x)$ can be defined in any $\mathcal{X} \subset \mathbb{R}^d$
- ▷ Practically, $\alpha_t(x)$ cannot be optimized using a grid optimizer.
- ▷ The size of the grid grows n^d and computational needs grow as $(n^d)^3$, where n number of grid points in 1D.
- ▷ One can use a first-order heuristic to optimize the acquisition function locally.
- ▷ More advanced methods: Look [Mutný and Krause, 2018] or visit:

Talk of Johannes tomorrow on BO for SwissFEL.

Part II, Programming: Lets try it out.

- Ekström, A., Forssén, C., Dimitrakakis, C., Dubhashi, D., Johansson, H., Muhammad, A., Salomonsson, H., and Schliep, A. (2019).
Bayesian optimization in ab initio nuclear physics.
arXiv preprint arXiv:1902.00941.
- Kirschner, J., Mutný, M., Hiller, N., Ischebeck, R., and Krause, A. (2019).
Adaptive and safe bayesian optimization in high dimensions via one-dimensional subspaces.
- McIntire, M., Cope, T., Ratner, D., and Ermon, S. (2016).
Bayesian optimization of fel performance at lcls.
Proceedings of IPAC2016.
- Mockus, J. (1982).
The bayesian approach to global optimization.
System Modeling and Optimization, pages 473–481.

-  Mutný, M. and Krause, A. (2018).
Efficient high dimensional bayesian optimization with additivity and quadrature fourier features.
In *Neural and Information Processing Systems (NeurIPS)*.
-  Rasmussen, C. E. (2004).
Gaussian processes in machine learning.
In *Advanced lectures on machine learning*, pages 63–71. Springer.
-  Romero, P. A., Krause, A., and Arnold, F. H. (2013).
Navigating the protein fitness landscape with gaussian processes.
Proceedings of the National Academy of Sciences (PNAS), 110(3).
-  Schneider, P.-I., Santiago, X. G., Soltwisch, V., Hammerschmidt, M., Burger, S., and Rockstuhl, C. (2018).

Benchmarking five global optimization approaches for nano-optical shape optimization and parameter reconstruction.